

Physics-inspired Neural Networks for building modeling and control

Loris Di Natale ^{1,2}

Bratislav Svetozarevic ¹

Philipp Heer ¹

Colin N. Jones ²

¹Urban Energy Systems Lab, Empa, Switzerland

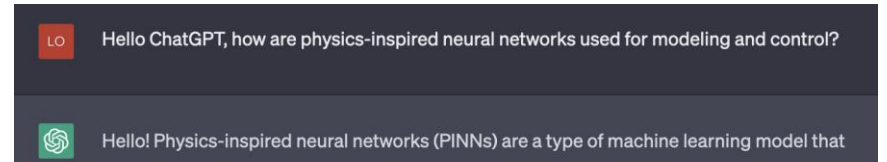
²Automatic Control Lab, EPFL, Switzerland

American Control Conference 2023, San Diego



Neural Networks and physical systems

- Neural Networks (NNs) achieve **amazing performance**...
- ... but can also **fail spectacularly**
 - They only do ***exactly*** what they are taught
- What about controlling (real-world) **physical systems with NNs?**
 - Data quantity/quality issues
 - Time to converge
 - Exploration phase
 - Guarantees?



Twitter taught Microsoft's AI chatbot to be a racist a*** in less than a day**

The Guardian




NEST, Empa, Dübendorf
Switzerland, © Zooey Braun

Go back to simulation

Neural Networks and physical systems

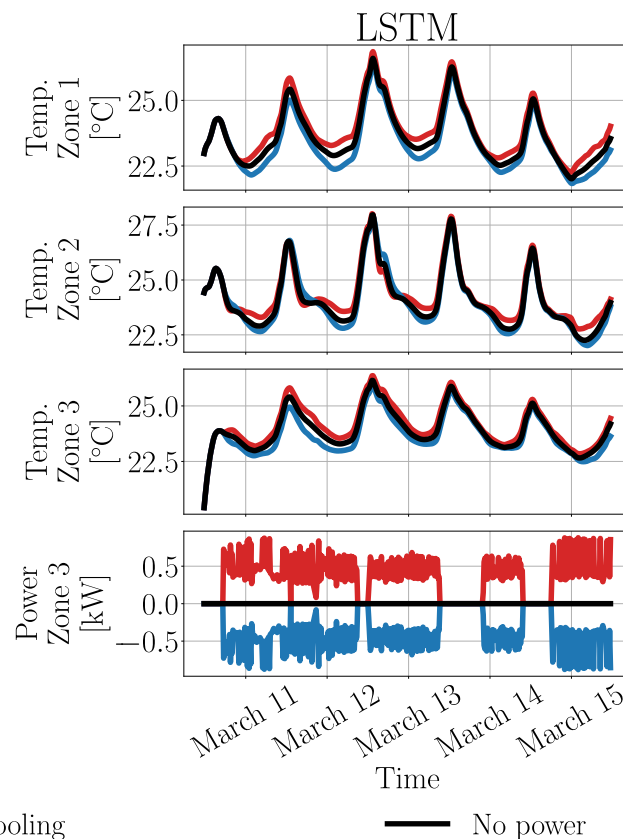
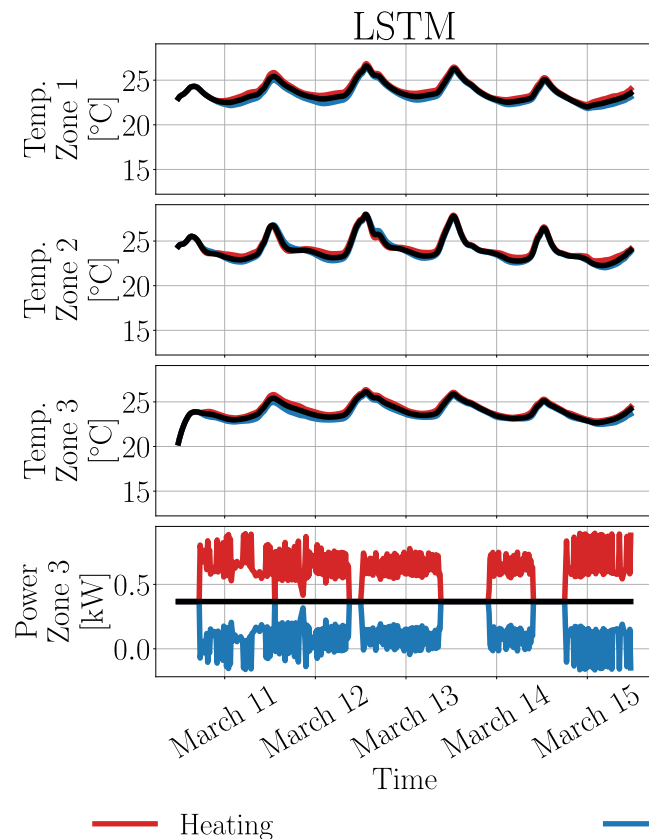

- Trained **sequentially** from data
- Can choose to use a **NN to model** the system as well



It works! Can save 25-30% of energy on UMAR compared to industrial baselines,^[1] but...

Neural Networks are physics-agnostic

- They attain **great accuracy**
- They can **violate physical laws** [2,3]





Heating and cooling have almost no (or even reversed) impact

*Can we help NNs?
We often have some prior knowledge of the system*

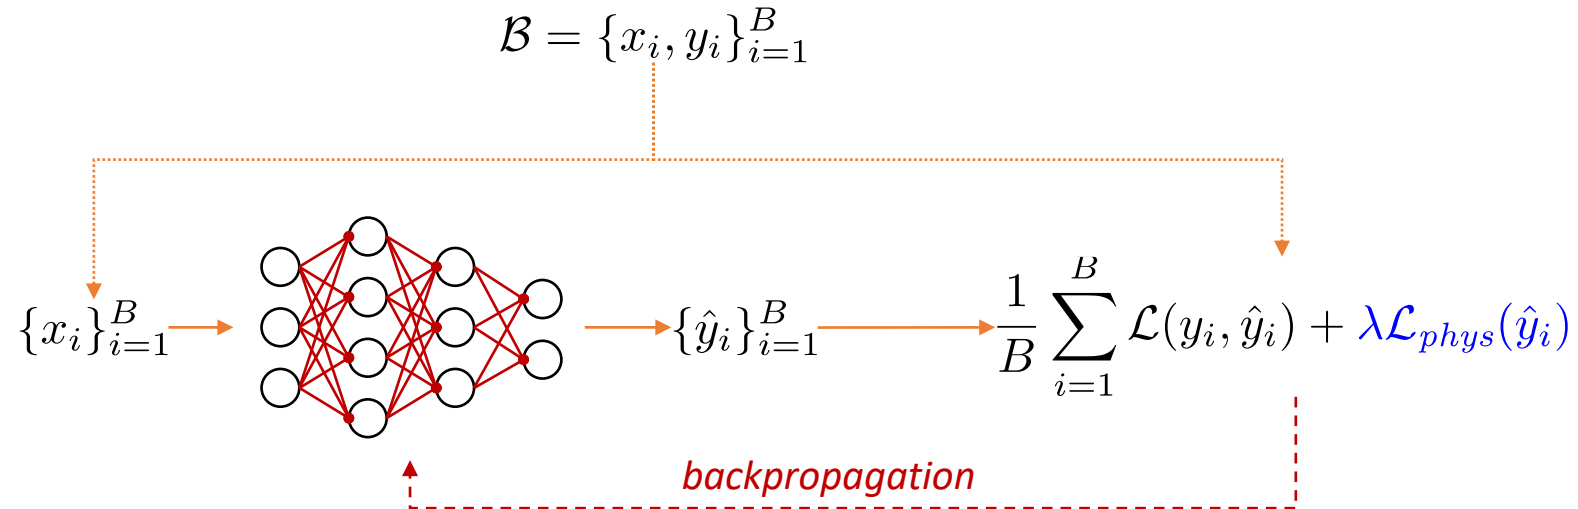

PiNNs for supervised learning

- Given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$



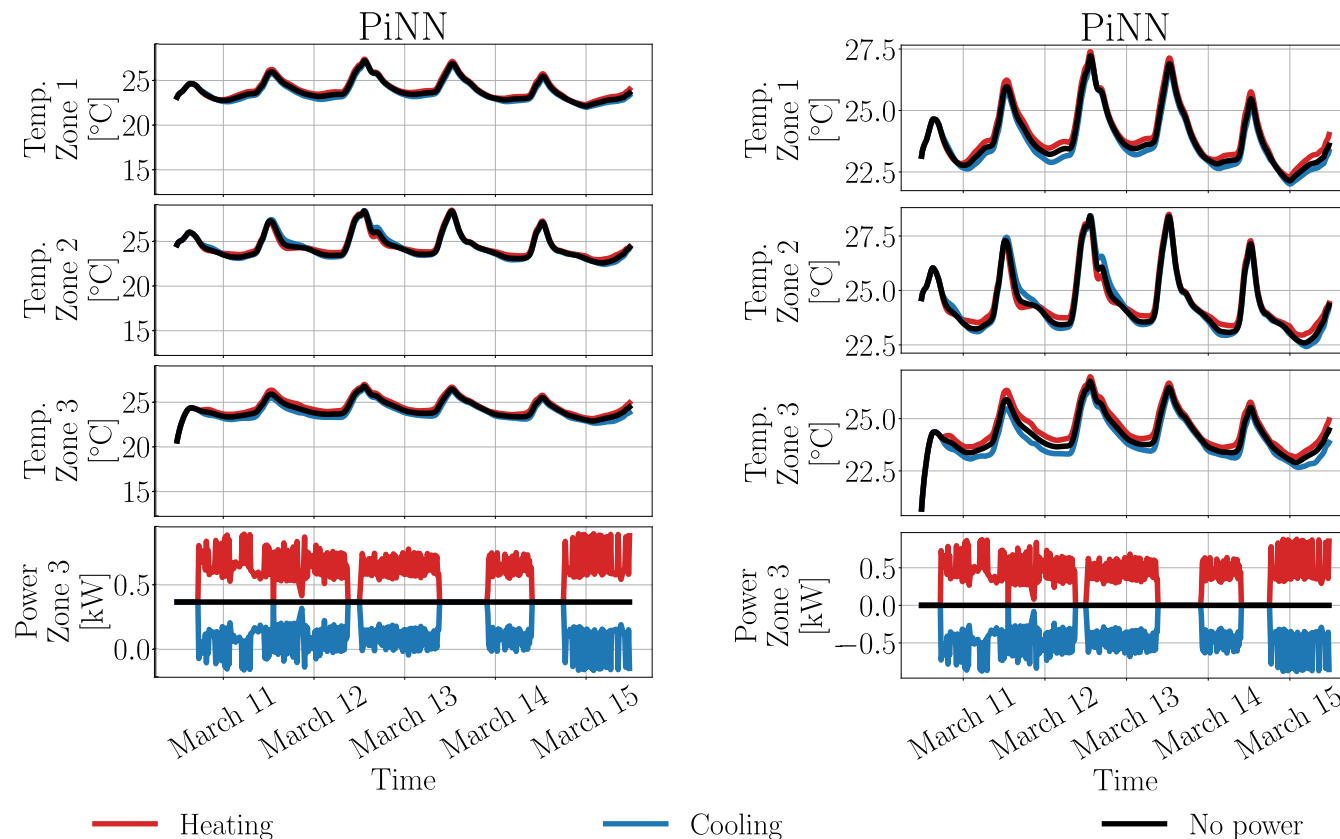

PiNNs for supervised learning

- Given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$
- Main idea: **Modified loss function** to steer the NN towards expected solutions
 - Boundary conditions
 - Physical laws (energy/mass/momentum conservation, ...)



Do PiNNs work?


- Want to ensure physical consistency: $\mathcal{L}_{phys}(\hat{y}_i) = \max \left(-\frac{\partial \hat{y}_i^{[a]}}{\partial x_i^{[b]}}, 0 \right)^{[3]}$
- Idea:** heating must heat the building



Does not solve our issues!

PiNNs for supervised learning

- Modified physical losses are **a good start to help NNs**



Pros

- Ease of implementation
- Can accelerate learning
- Can improve the solution

$$\frac{1}{B} \sum_{i=1}^B \mathcal{L}(y_i, \hat{y}_i) + \lambda \mathcal{L}_{phys}(\hat{z}_i, z_i)$$

Cons


- Hyperparameter tuning
- **No guarantees**

*Whenever possible, you can also try to use **tailored** architectures*

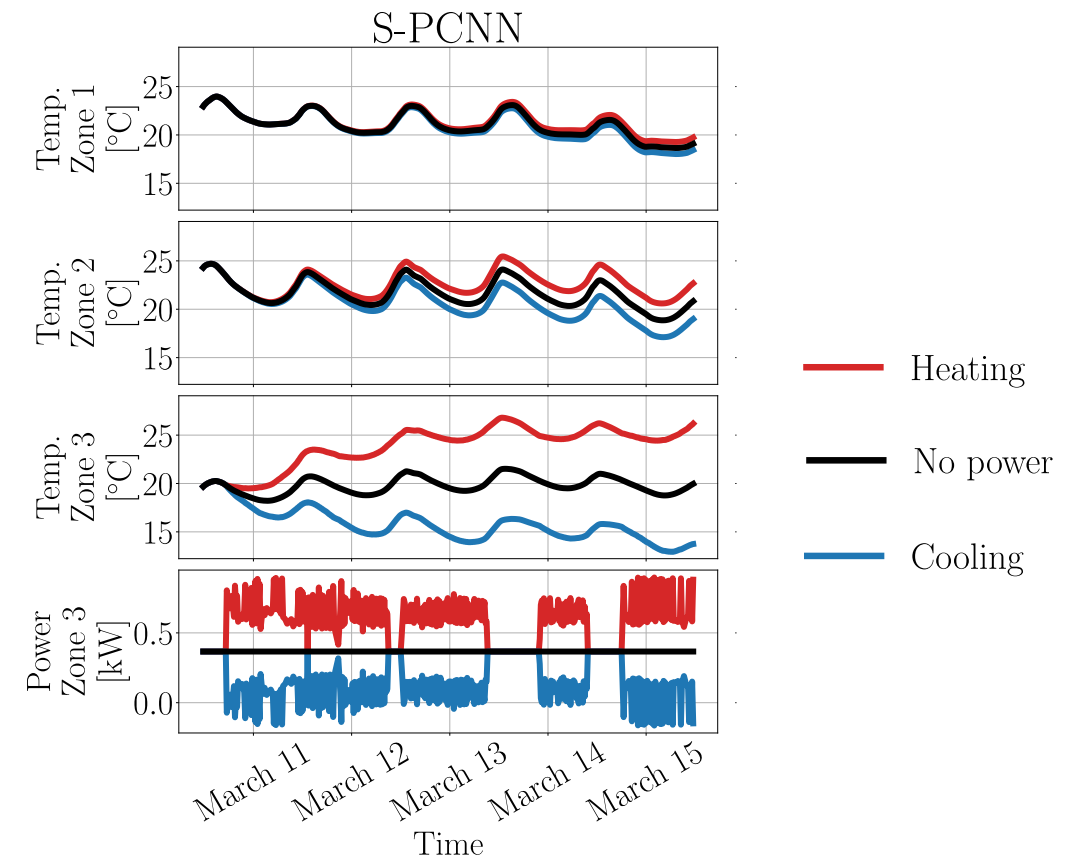

- Hamiltonian NNs,^[4] Lagrangian NNs,^[5] ...

Physically Consistent NNs (PCNNs) [6]


- Let a basic **physics-inspired module** run in parallel of the NN



- Key:** Train θ **simultaneously** through backpropagation
- State-of-the-art accuracy** on the analyzed case study
- But **following the laws of physics**





Neural Networks and physical systems



*If you have a good model, you can use **MPC**, unless it is impractical*

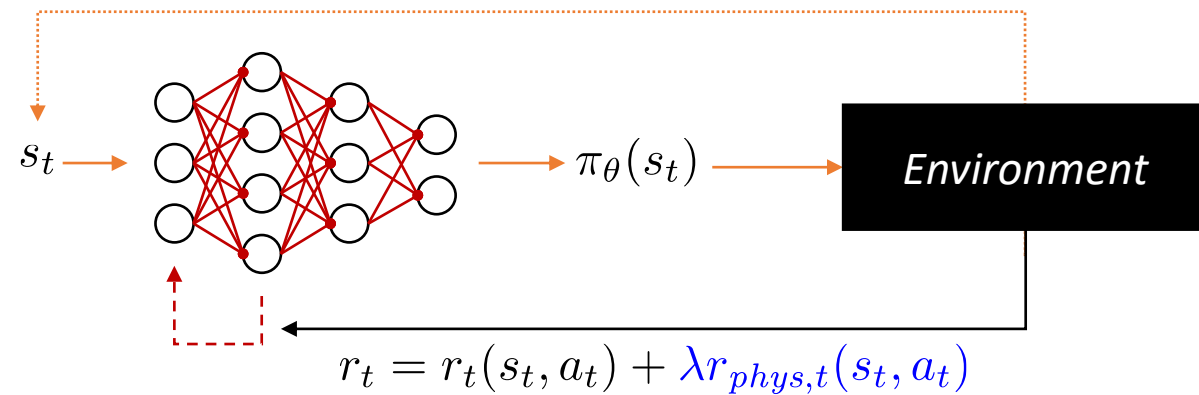

PiNNs for reinforcement learning

- Learning through interaction with the environment





PiNNs for reinforcement learning

- Learning through interaction with the environment
- **Reward shaping**: modify the feedback to the network (e.g., constraints violations, ...)



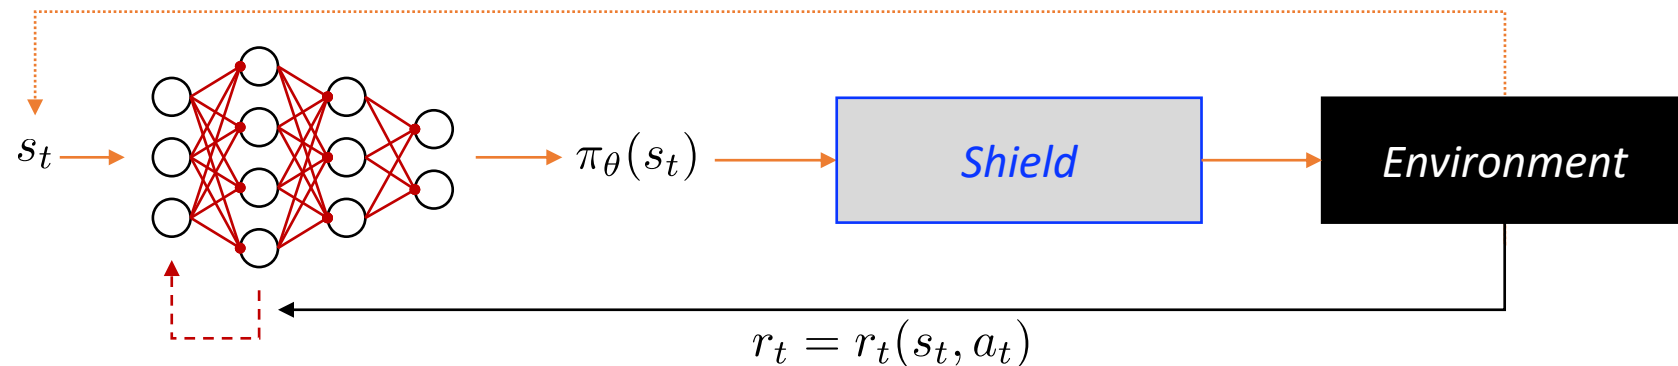

PiNNs for reinforcement learning

- Learning through interaction with the environment
- **Reward shaping:** modify the feedback to the network (e.g., constraints violations, ...)
- **Residual learning:** additive, multiplicative, time varying, ...



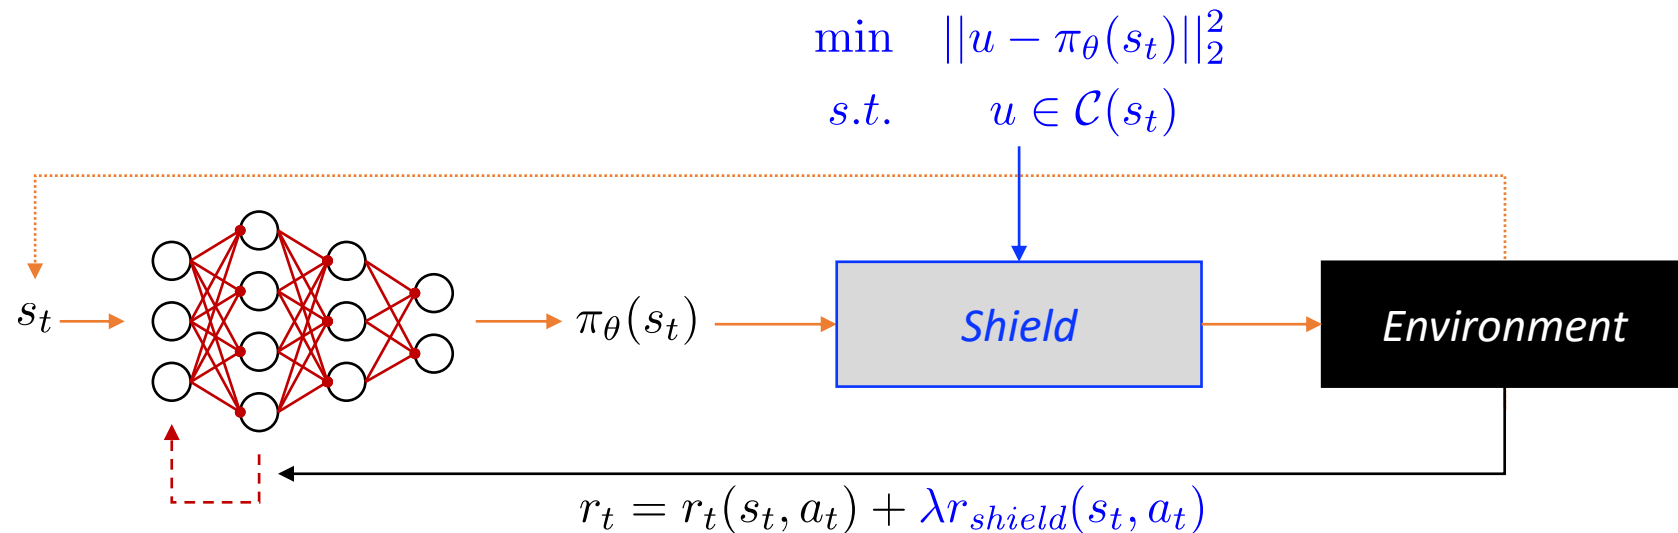

PiNNs for reinforcement learning

- Learning through interaction with the environment
- **Reward shaping:** modify the feedback to the network (e.g., constraints violations, ...)
- **Residual learning:** additive, multiplicative, time varying, ...
- **Shielding:** Define a filter $\mathcal{C}(s_t)$ to correct the agent when some specifications are not met



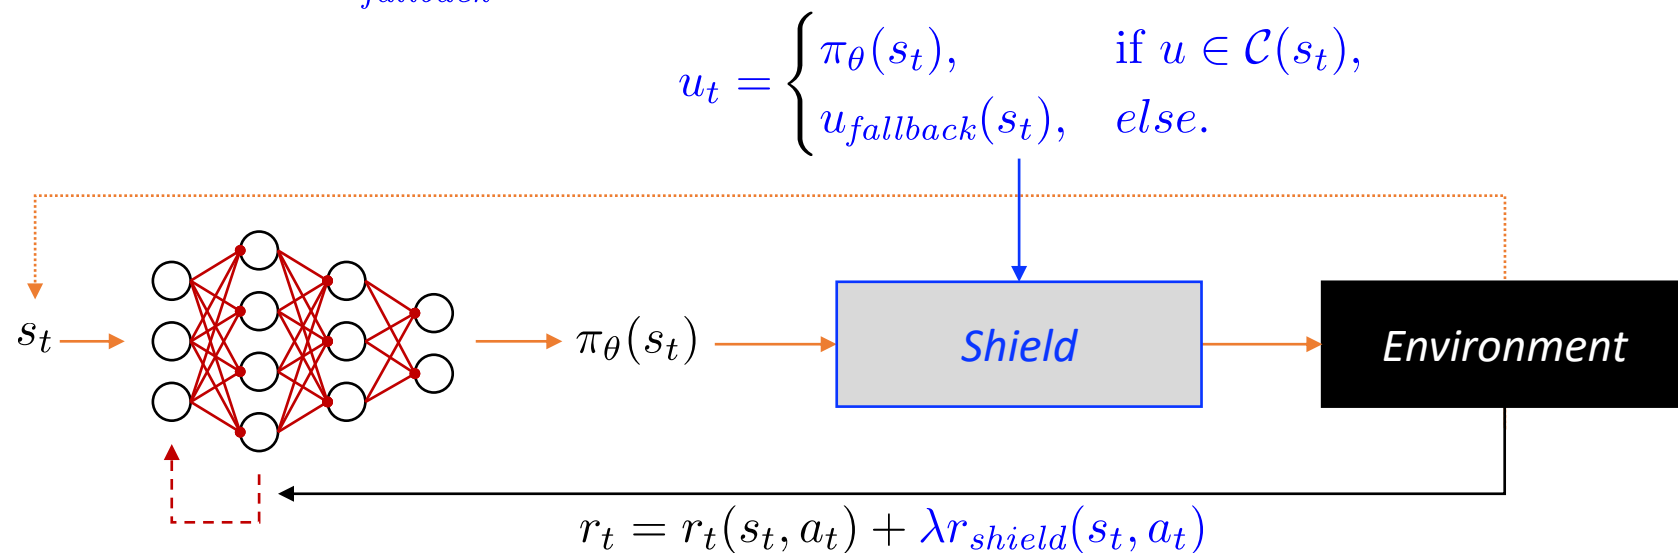

PiNNs for reinforcement learning

- Learning through interaction with the environment
- **Reward shaping**: modify the feedback to the network (e.g., constraints violations, ...)
- **Residual learning**: additive, multiplicative, time varying, ...
- **Shielding**: Define a filter $\mathcal{C}(s_t)$ to correct the agent when some specifications are not met
 - Project on the desired set




PiNNs for reinforcement learning

- Learning through interaction with the environment
- **Reward shaping**: modify the feedback to the network (e.g., constraints violations, ...)
- **Residual learning**: additive, multiplicative, time varying, ...
- **Shielding**: Define a filter $\mathcal{C}(s_t)$ to correct the agent when some specifications are not met
 - Project on the desired set
 - Use a known fallback controller $u_{fallback}$



PiNNs for reinforcement learning

- Introducing physics/knowledge can **accelerate learning** and **lead to better solutions**




*Can we improve the computational efficiency of **shielding** to enforce prior knowledge?*

Computationally Efficient RL [7]

- Use time-varying bounds on the agent's actions to **avoid suboptimal state-action pairs**

"If it's too cold, then heat"
"If it's too hot, then don't heat"
...

$$\longrightarrow \mathcal{C}(s_t) = u^{\min}(s_t) \leq u \leq u^{\max}(s_t)$$





- The projection is **computationally cheap**: saturate the actions at each step!
- Modify the gradient update** of the agent so it learns from its mistakes

*No computational overhead
but no guarantees*

Computationally Efficient RL [7]



- Use time-varying bounds on the agent's actions to **avoid suboptimal state-action pairs**



*No computational overhead
but no guarantees*

Computationally Efficient RL [7]



- Use time-varying bounds on the agent's actions to **avoid suboptimal state-action pairs**



*No computational overhead
but no guarantees*

Computationally Efficient RL [7]



- Use time-varying bounds on the agent's actions to **avoid suboptimal state-action pairs**



*No computational overhead
but no guarantees*

Computationally Efficient RL [7]



- Use time-varying bounds on the agent's actions to **avoid suboptimal state-action pairs**



*No computational overhead
but no guarantees*

Computationally Efficient RL [7]

- Use time-varying bounds on the agent's actions to **avoid suboptimal state-action pairs**



No computational overhead
but no guarantees

Up to **6-7x** better sample
complexity than classical DRL

Up to **2-3x** better sample
complexity than RS

Key takeaways

- Neural Networks are powerful but **physics-agnostic**

PiNNs for modeling

You can try different **informative loss functions**

If possible, you can try **tailored architectures** like PCNNs

PiNNs for control

You can try different **reward functions** and **residual learning techniques**

Shielding can give guarantees but is **usually complex**



Thank you for your attention!

Loris DI NATALE


Urban Energy Systems Lab, Empa
Laboratoire d'Automatique, EPFL

loris.dinatale@empa.ch
<https://www.empa.ch/web/dilo>

References

- [1] Di Natale, L., Svetozarevic, B., Heer, P., & Jones, C. N. (2022, June). Near-optimal deep reinforcement learning policies from data for zone temperature control. In *2022 IEEE 17th International Conference on Control & Automation (ICCA)* (pp. 698-703). IEEE.
- [2] Geirhos, R., Jacobsen, J. H., Michaelis, C., Zemel, R., Brendel, W., Bethge, M., & Wichmann, F. A. (2020). Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11), 665-673.
- [3] Di Natale, L., Svetozarevic, B., Heer, P., & Jones, C. N. (2023). Towards scalable physically consistent neural networks: An application to data-driven multi-zone thermal building models. *Applied Energy*, 340, 121071.
- [4] Greydanus, S., Dzamba, M., & Yosinski, J. (2019). Hamiltonian neural networks. *Advances in neural information processing systems*, 32.
- [5] Cranmer, M., Greydanus, S., Hoyer, S., Battaglia, P., Spergel, D., & Ho, S. (2020). Lagrangian neural networks. *arXiv preprint arXiv:2003.04630*.
- [6] Di Natale, L., Svetozarevic, B., Heer, P., & Jones, C. N. (2022). Physically consistent neural networks for building thermal modeling: theory and analysis. *Applied Energy*, 325, 119806.
- [7] Di Natale, L., Svetozarevic, B., Heer, P., & Jones, C. N. (2022). Efficient Reinforcement Learning (ERL): Targeted Exploration Through Action Saturation. *arXiv preprint arXiv:2211.16691*.

Appendix



Key takeaways

- Neural Networks are powerful but **physics-agnostic**

PiNNs for modeling

You can try different **informative loss functions**

If possible, you can try **tailored architectures** like PCNNs

PiNNs for control

You can try different **reward functions** and **residual learning techniques**

Shielding can give guarantees but is **usually complex**

- Conclusion


Key takeaways

- Neural Networks are powerful but **physics-agnostic – but solutions exist!**
- **PiNNs for modeling** – supervised learning
 - Implementing a physical loss can be **easy and might help** but tuning the weighting factor is not trivial
 - Use **tailored architectures** when possible
 - PCNNs for example attain **state-of-the-art accuracy** for building thermal modeling **while ensuring physical consistency**
- **PiNNs for control** – reinforcement learning
 - While **reward shaping** and **residual learning** are easy, they might not enforce knowledge
 - **Shielding** can preserve prior specifications but is often computationally expensive
 - **Enforcing simple rules instead can already help without computational overhead**


Use (a combination of) these techniques to improve your NNs

- Physics-inspired NNs (PiNNs)

PiNNs for supervised learning




- Given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$
- Main idea:** Modified **loss function** to **steer** the NN towards expected solutions
 - Boundary conditions
 - Physical laws (energy/mass conservation, ...)
- Can **augment the NN's output**

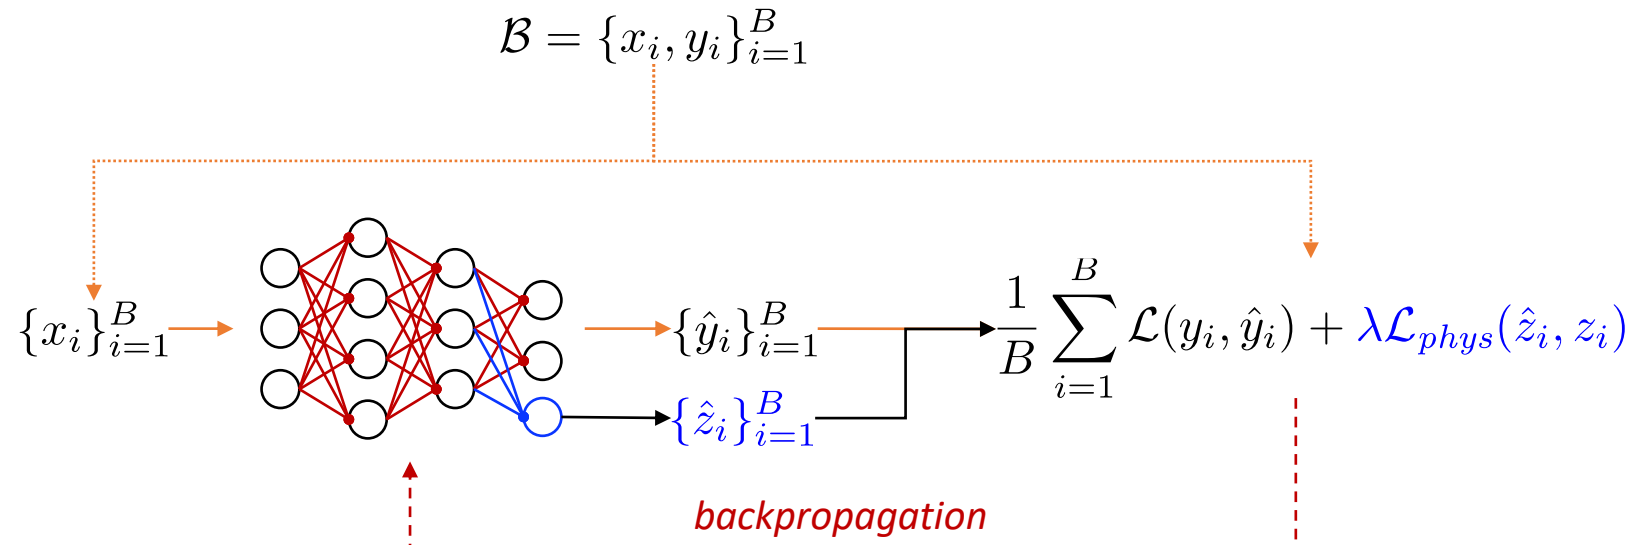


- Physics-inspired NNs (PiNNs)

PiNNs for supervised learning




- Given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$
- Main idea:** Modified **loss function** to **steer** the NN towards expected solutions
 - Boundary conditions
 - Physical laws (energy/mass conservation, ...)
- Can **augment the NN's output**

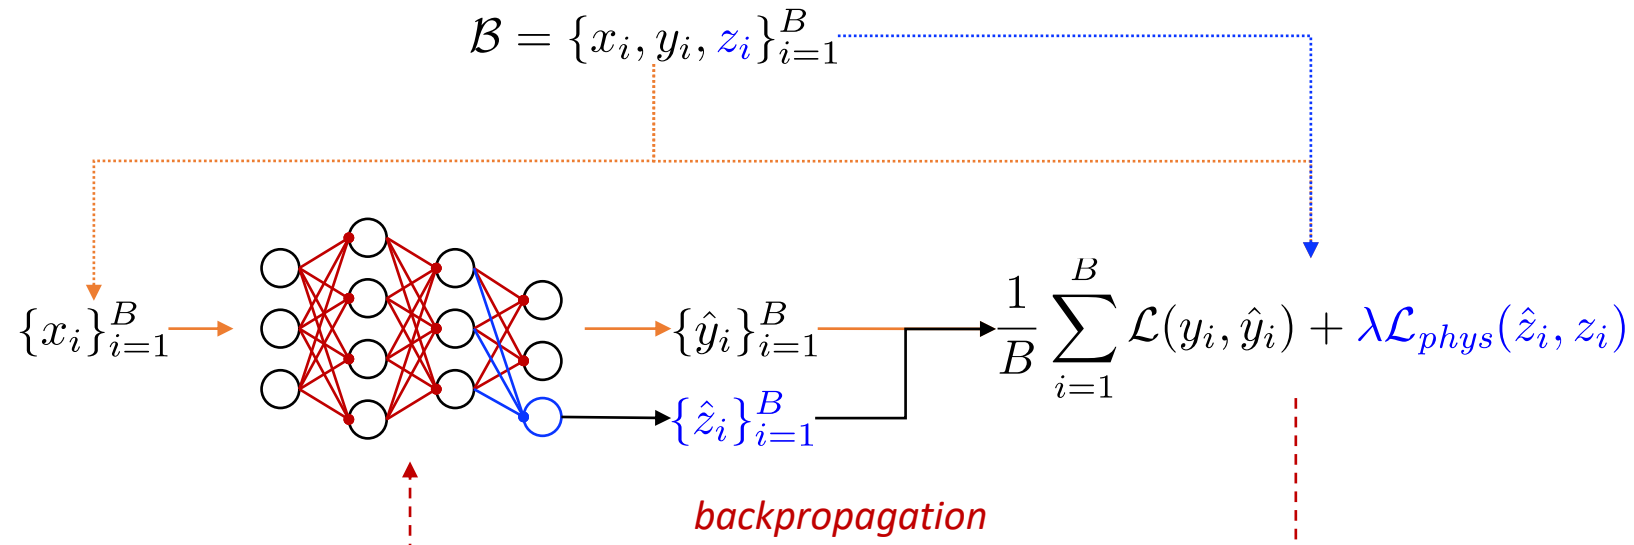


- Physics-inspired NNs (PiNNs)

PiNNs for supervised learning




- Given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$
- Main idea:** Modified **loss function** to **steer** the NN towards expected solutions
 - Boundary conditions
 - Physical laws (energy/mass conservation, ...)
- Can **augment the NN's output**

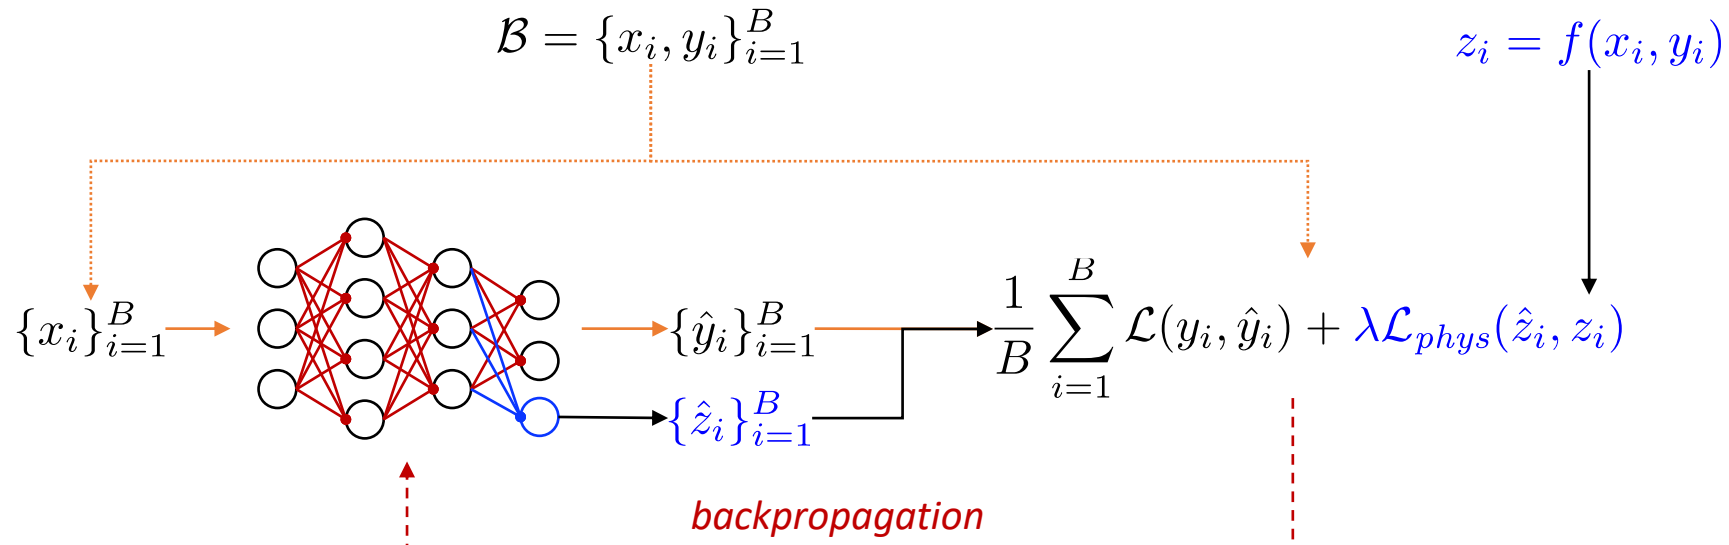


- Physics-inspired NNs (PiNNs)

PiNNs for supervised learning




- Given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$
- Main idea:** Modified **loss function** to **steer** the NN towards expected solutions
 - Boundary conditions
 - Physical laws (energy/mass conservation, ...)
- Can **augment the NN's output**
- Can **create a bottleneck**




- Physics-inspired NNs (PiNNs)

PiNNs for supervised learning



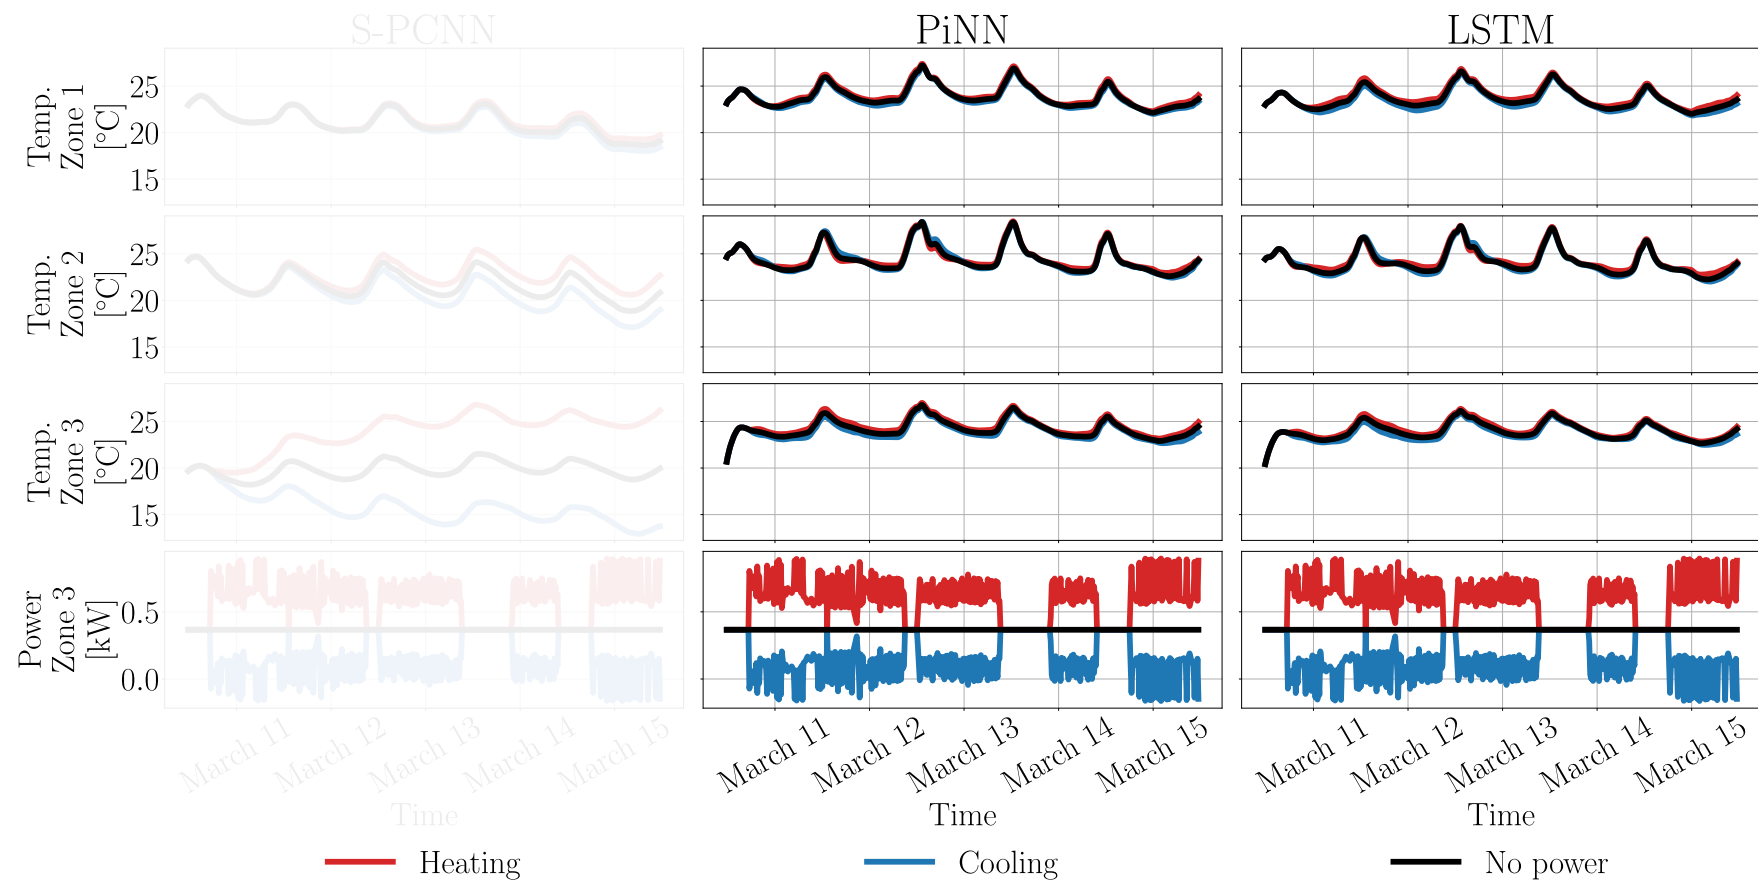
- Given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$
- Main idea:** Modified **loss function** to **steer** the NN towards expected solutions
 - Boundary conditions
 - Physical laws (energy/mass conservation, ...)
- Can **augment the NN's output**
- Can **create a bottleneck**
- ...



The issue of *shortcut* learning

- Case study on UMAR: learn the **thermodynamics of buildings**
- PiNNs and LSTMs have a **great accuracy** on the **validation data**

Issue of *shortcut learning*
No impact of heating/cooling!



Physically Consistent NNs (PCNNs)

- Enforce physical properties *by design* in a module running in parallel of the NN


Building thermodynamics example

We want to ensure:

- The predicted temperatures increase when the applied heating power or the outside temperature increases;
- Physically meaningful energy flows between the thermal zones.

This is **hard-encoded in the physics-inspired module**, which processes the power inputs, ambient temperatures, and energy exchanges between the zones.

- The parameters are **learned from data**




$$\begin{aligned}
 D_{k+1} &= D_k + f(x_k, D_k), \\
 E_{k+1} &= E_k + a_h \max\{u_k, 0\} + a_c \min\{u_k, 0\} \\
 &\quad - b(T_k - T_k^{out}) - \sum_{z' \in \mathcal{N}(z)} c_{z'}(T_k - T_k^{z'}), \\
 T_{k+1} &= D_{k+1} + E_{k+1},
 \end{aligned}
 \quad \left\{ \begin{array}{l} \frac{\partial T_{k+i}^z}{\partial u_{k+j}^z} > 0 \\ \frac{\partial T_{k+i}^z}{\partial T_{k+j}^{out}} > 0 \\ \frac{\partial T_{k+i}^z}{\partial T_{k+j}^y} > 0 \end{array} \right. *$$

* under mild assumptions on a_h, a_c, b, c


Solving the *shortcut* learning issue

- Case study on UMAR: learn the **thermodynamics of buildings**
- PiNNs and LSTMs have a **great accuracy** on the **validation data** – **PCNNs are physically consistent by construction**




Solving the *shortcut* learning issue

- Case study on UMAR: learn the **thermodynamics of buildings**
- PiNNs and LSTMs have a **great accuracy** on the **validation data** – PCNNs are **physically consistent *by construction***
- But they also **outperform all the other physically consistent methods by 17-35%**



Simulation

- DRL agents **converge to behaviors similar to the optimal trajectories** (computed *a posteriori*)
 - Confirmed by **statistical analysis** over 2000 3-day long sequences
 - And for **different reward functions** and **random seeds**



Experiment

- Only short period because of connection issues
- But the agent shows **similar behavior as in simulation**
 - Preheating
 - Close to lower bound




- Introduction

An issue of convergence speed

- **3 years** worth of data (randomly sampled days) to obtain the performance of classical rule-based controllers
- **15 years (500k points)** to “converge”
- (D)RL is usually (*very, very, very*) **data-inefficient**

High computational burden

Limited deployment on physical systems



- Introduction


Why?

To get there...

... we need this!

Exploration is *necessary*

... but often *stupid!*



Postulate

Prior expert knowledge often provides **intuition** about what optimal policies should do


We can accordingly **constrain the agents** to regions of the state space that are deemed interesting

- Methods

Proposed approach to *accelerate learning*

How can we constrain the actions of agents to explore interesting states?

- Design **handcrafted simple rules** from expert knowledge to ensure meaningful actions
- **Saturate** the agents' actions according to these rules
- **Modify the gradient update** of the policy to let it learn from the corrections and steer it towards expected behaviors



- Methods

Proposed approach to *accelerate learning*

How can we constrain the actions of agents to explore interesting states?

- Correct the agents using safe action sets - **but no need to always satisfy these artificial constraints!**

- Design **handcrafted simple rules** from expert knowledge to ensure meaningful actions
- **Saturate** the agents' actions according to these rules
- **Modify the gradient update** of the policy to let it learn from the corrections and steer it towards expected behaviors

$$a^{min}(s), a^{max}(s) = f(s)$$

$$a(s) = \text{clip}(\pi_{\theta}(s) + \nu, a^{min}(s), a^{max}(s)) \quad \nu \sim \mathcal{N}(0, \sigma^2)$$


$$\hat{\nabla}_{\theta}^{EA} \pi_{\theta} = \nabla_{\theta} \left[\frac{1}{|B|} \sum_{s \in B} \frac{1}{|B|} \sum_{a \in B} Q_{\phi}(\mathcal{Q}_{\phi}(\pi_{\theta}(s), a)(s)) \right] - \frac{\lambda}{2} (\pi_{\theta}(s) - a(s))^2$$

Each step is easy to implement and computationally inexpensive

- Introduction

What does the saturation do?

*To get there...
... we need this!*





— Classical DRL — EA 0 / 1 — EA 0 / 0.5 — EA 0 / 0.25

Meaningful exploration

- Results

Does it impact performance?


- The learnt behaviors are similar!



Good final performance


- Results

Does it accelerate learning?




RS 0 / 1

EA 0 / 1



EA 0.5 / 1

EA 0 / 1




Efficient Agents (EAs) (ours)

Classical Agents



Impact of m on EAs (ours) – more freedom



RS 0 / 0.25

EA 0 / 0.25



RS 0 / 0.1

EA 0 / 0.1


Smaller n
tighter constraints

Main metric: number of *days of data* required to converged to the performance of industrial baselines

0.1


- Results

Does it accelerate learning?



- Results

Does it accelerate learning?



Objective attained! Up to 6-7x faster

- Wrap-up

Discussion & Conclusion

- **Black-box models can be misleading**, even if they fit the data well
 - PCNNs as a **potential solution - easy to scale**
- **Deep Reinforcement Learning** provides good results when coupled with PCNNs
- It is possible to design **simple rules based on prior knowledge** to help (D)RL agents converge faster
 - **Saturate** the actions of the agents accordingly
 - **Modify the gradient update** to let agents learn from it

→ *Avoid tedious engineering*

→ *Energy savings of 25-30%*

→ *Comfort of the occupants improved*

→ *Straightforward implementation*

→ *Computationally inexpensive*

→ *Up to 6-7 times faster*

→ Fully **black-box pipeline** from data to control policies

→ *“Plug & Play” controllers*

- Wrap-up

Discussion & Conclusion

- **Neural Networks can be misleading**, even if they fit the data well
 - PCNNs as a **potential solution** - **easy to scale**
 - **Deep Reinforcement Learning** provides good results when coupled with PCNNs
 - It is possible to design **simple rules based on prior knowledge** to help (D)RL agents converge faster
 - **Saturate** the actions of the agents accordingly
 - **Modify the gradient update** to let agents learn from it
- Fully **black-box pipeline** from data to control policies
- *State-of-the-art accuracy*
→ *Physical consistency*
- *Energy savings of 25-30%*
→ *Comfort of the occupants improved*
- *Straightforward implementation*
→ *Computationally inexpensive*
→ *Up to 6-7 times faster*
- *“Plug & Play” controllers*

Main references

PCNNs

- **Towards Scalable Physically Consistent Neural Networks: an Application to Data-driven Multi-zone Thermal Building Models.**
L. Di Natale, B. Svetozarevic, P. Heer, and C.N. Jones (2022).
Manuscript submitted to Applied Energy. <https://arxiv.org/abs/2212.12380>.
- **Physically consistent neural network building models: theory and analysis.**
L. Di Natale, B. Svetozarevic, P. Heer, and C.N. Jones (2022).
Applied Energy, 119806. <https://doi.org/10.1016/j.apenergy.2022.119806>.

DRL

- **Efficient Reinforcement Learning (ERL): Targeted Exploration Through Action Saturation.**
L. Di Natale, B. Svetozarevic, P. Heer, and C.N. Jones (2022).
Manuscript submitted to L4DC 2023. <https://arxiv.org/abs/2211.16691>.
- **Near-optimal Deep Reinforcement Learning Policies from Data for Zone Temperature Control.**
L. Di Natale, B. Svetozarevic, P. Heer, and C.N. Jones (2022).
2022 IEEE 17th International Conference on Control & Automation (ICCA), 698-703.
<https://doi.org/10.1109/ICCA54724.2022.9831914>.